

Kubernetes

YAML

- Kind deployment or kind service
- ETCD holds the cluster status
- Deployment has configuration and then also pod configuration
- Key value pairs for components app: nginx
- Targetport should match the containerport, else they cant communicate with each other
- kubectl get pod -o wide for more information

Deployment

- Replica's specified the number of pods at any given time
- Selector identifies which pods belong to this deployment
- Labels are attached to the deployment itself
- The selector must match the labels applied to the deployment itself
- Starts at spec

Pod

- Starts at template

Containers

- Starts at containers

Service

- Selector has a name that references the pod so in the mongo db example it was mongodb
- To expose the service to the outside use LoadBalancer type and add nodeport 30000-32767

Deployment is the what, the service the how. They rarely work without each other

Namespace

- kubectl create namespace <bla>
- kubectl create ns

Why do I need a namespace?

- Separation
 - Database

- Monitoring
- Elastic stack
- Nginx ingress
- Multiple teams in a kubernetes cluster
 - Only access to their own cluster

—

Change the namespace instead of default

- `kubectl config set-context —current namespace=<namespace>`
-

Services

Stable IP adress

Loadbalancing

Loose coupling

Within & Outside cluster

ClusterIP service

- Gives an Ip range
- The third number in the IP specifies the node which is on
- Service registers endpoints in the deployment app: my-app to each other, so it knows they belong to each other.
- Service will find the pods based on labels and selector and send the request based on targetport
- Service target port needs to match the containerport van the pod
- If you have multiple ports, you have to name those ports.

Headless service

- Most commonly used with databases (Mysql, postgres or mongo)
- **statefulset**
- Pods are not identical
- DNS lookup for a service
- clusterIP None

Nodeport service

- Predefined values between 30000 to 32000
- External traffic
- Testing service
- Not for production

Loadbalancer service

- Server is accessible from outside

StatefulSet

- Stateful applications
 - That track state all databases
 - Mysql main, and 2 replicas if you have 3 database pods. Data is else not interchangeable
 - They dont have access to the same storage
 - Use data persistence storage
 - Own DNS
- Stateless applications
 - Dont keep records of state

HELM

- `chmod 400 k8s-test.yaml` allows only your user to access the kube config
- `export KUBECONFIG=`
- `Helm install <name we give> --values helm-mongodb.yaml bitnami/mongodb`

CONFIG MAP & Secret volume

- Config map & Secret are used for external configuration
- Adding the volume files first into the pod, but that does not mean its in the container. The volumeMount needs to be in the container level
- VolumeMount name is the name of the volume that is attached to the POD.
- `kubectl get configmap prometheus-config -n default -o yaml`
- Left part of docker compose (host path is volume in Kubernetes)
- Right part of docker compose is volumeMount and specifies the container

Volumes

- **Persistent volume**
 - Need physical storage
 - Available to the whole cluster
 - Local
 - Tied to one specific node
 - Surviving cluster crashes
 - Remote volume type
 - DB persistence should use remote storage
- K8s admin
 - Sets up and maintains the cluster
- K8s user
 - Deploys applications in cluster
 - Explicitly set the Persistent claim

- Pod needs to use PV claim and claim name needs to be same as the Persistent Volume claim yaml
 - Left part of docker compose (host path is volume in Kubernetes)
 - Right part of docker compose is volumeMount and specifies the container (right)
 - subpath: Maps the file, if you would not specify the submap it would take the entire directory
 - **Persistent volume claim**
-
- **Storage class**
 - Uses yaml as well
 - provisioner is needed in the yaml
 - Requested by PVC, storage classname needs to be defined in the yaml in order to specify

Kubernetes operators

- **Stateless applications**
 - k8s can manage this in automated way
- **Stateful applications**
 - Should be updated and destroyed in order
 - Need manual intervention
 - **Operator**
 - Deployment
 - service
 - configmap
 - Statefulset
 - **CRD**
- Operatorhub

RBAC

- Permissions for a cluster
- Role only defines resources and access permissions
- Role binding (Link role to a user or user group)
- Cluster role defines what resources have what permissions, clusterwide.
- External sources for user management
 - Admin will configure external sources
- Application only get permissions to work
- Service account component inside the cluster as a representation of that user
- Rules
 - API groups
 - Resources (pods)

- verbs (you can add crud functionality here)
- Resourcenames (here you can specify in example the app or the db and specific permissions for those items)

Kubectl commands

- kubectl get pods (gets the pods)
- kubectl get services (gets services)
- kubectl create deployment nginx-depl --image=nginx
- kubectl get replicaset
- kubectl logs <podname> this will log
- kubectl describe pod mongo-deployment-555654868f-xn2q7
- kubectl exec -it <podname> — bin/bash
- Deployment > Pod > Container
- kubectl get all | grep mongodb
-

Ingress

- Ingress controller is needed to be able to use ingress and is separated from the release
- TLS certificate needs to be hosted in the cluster to give it https
- Only values no files
- Single endpoint into the cluster